



ELTE | FACULTY OF
INFORMATICS

Supervised Learning

Introduction to Machine Learning - Lecture 4

Computer Science BSc Course, ELTE Faculty of Informatics

János Botzheim

Associate Professor

Department of Artificial Intelligence

botzheim@inf.elte.hu



DEPARTMENT OF
ARTIFICIAL
INTELLIGENCE

Unsupervised learning



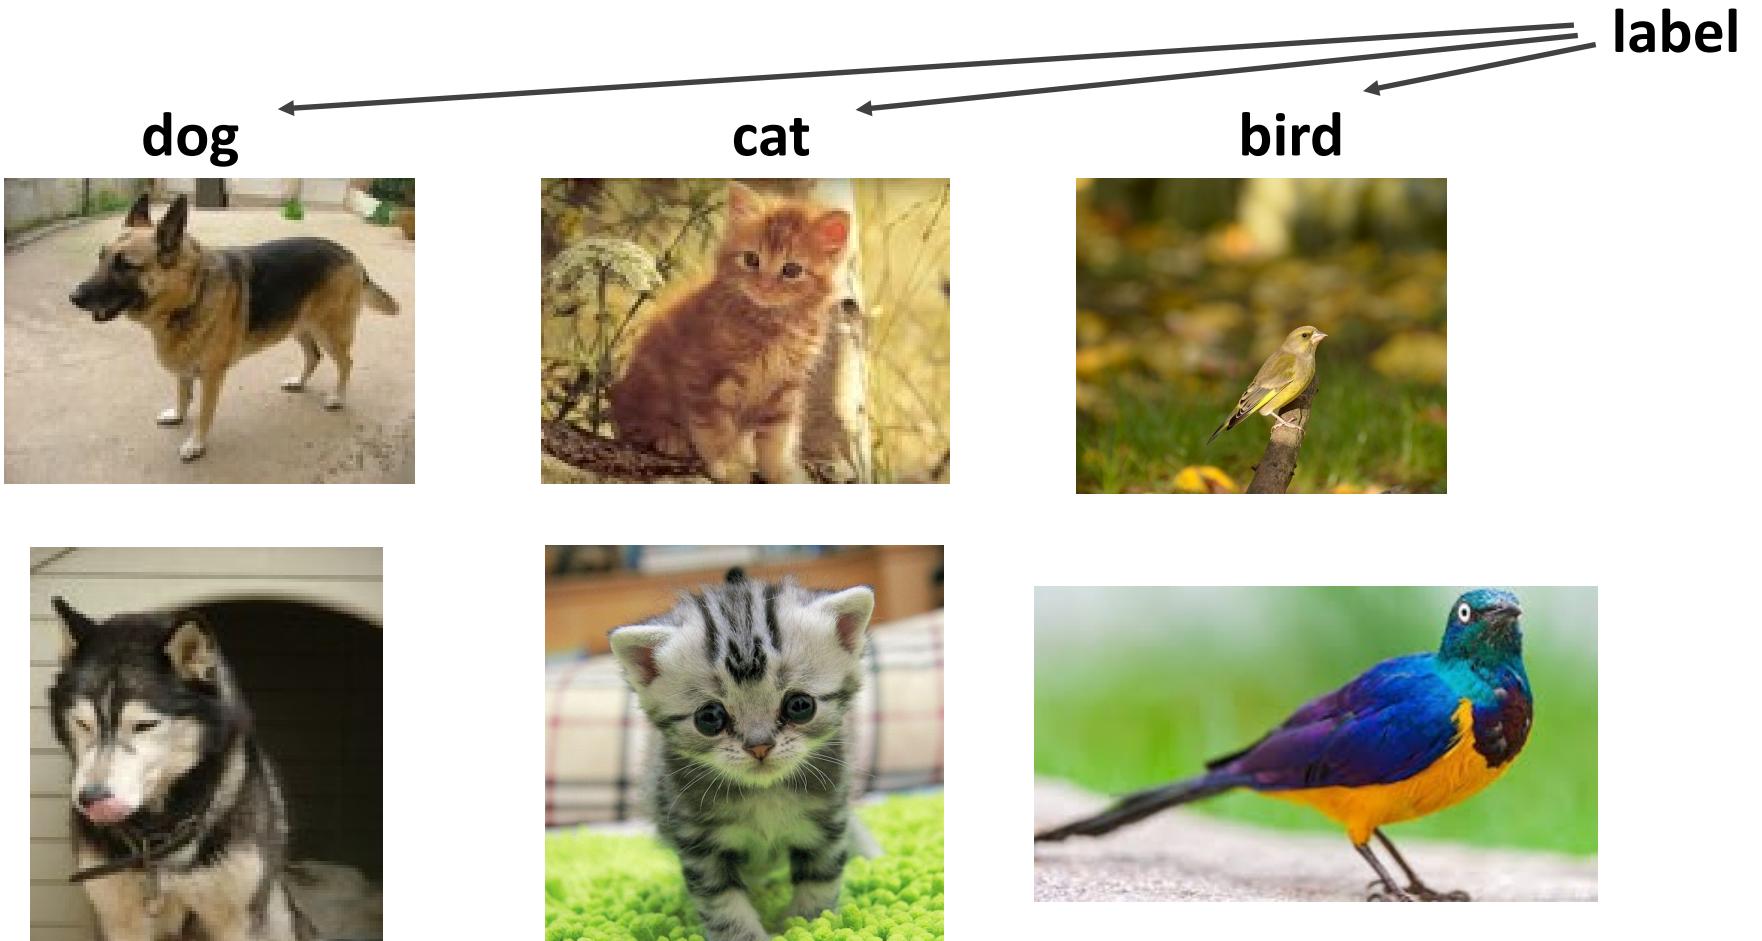
ELTE

FACULTY OF
INFORMATICS

Supervised learning

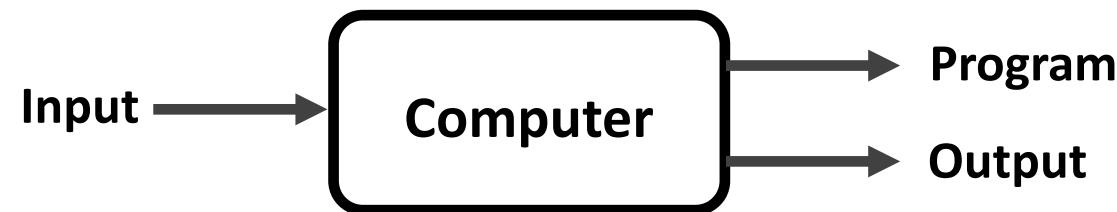
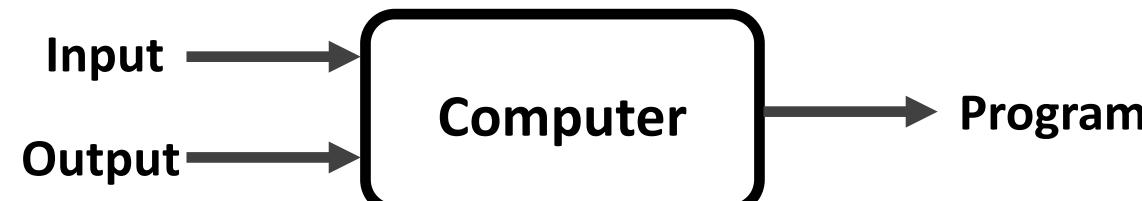
Introduction

Supervised learning



Learning - Programming

- Traditional computer programming:
- Supervised learning:
- Unsupervised learning:



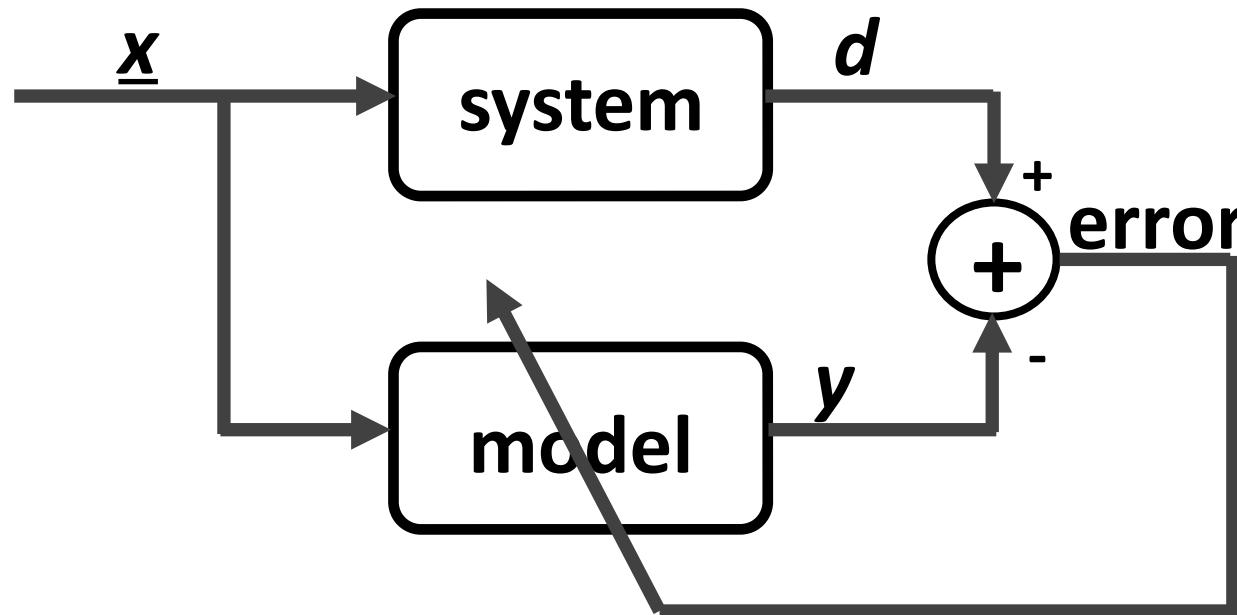
Types of learning

- **Learning**
 - Agent can improve its performance based on observations
- **Machine Learning**
 - Subset of artificial intelligence, computer system learns by patterns in observation data
- Learning tasks
 - **Classification:** output is one of a finite set of values
 - **Regression:** output is a numeric prediction
- Learning Feedback
 - **Supervised Learning** – learn a function that maps inputs to outputs by observing input-output pairs
 - **Unsupervised Learning** – learn patterns in data without explicit feedback (e.g. clustering)
 - **Reinforcement Learning** – learn beneficial actions based on rewards and punishments



Supervised learning

- given: input-output training patterns $(x_1^{(p)}, x_2^{(p)}, \dots, x_n^{(p)}, d^{(p)})$
- p : number of patterns, n -dimensional input, scalar output



Supervised learning

- **Representation:**
 - rule based system
 - fuzzy system
 - decision tree
 - neural network
 - support vector machine
 - etc.
- **Evaluation:**
 - error
 - accuracy
 - cost
 - entropy
 - etc.
- **Optimization:**
 - gradient based algorithm
 - evolutionary algorithm
 - etc.



Supervised Learning

- More formally

Given a **training set** of N example input–output pairs

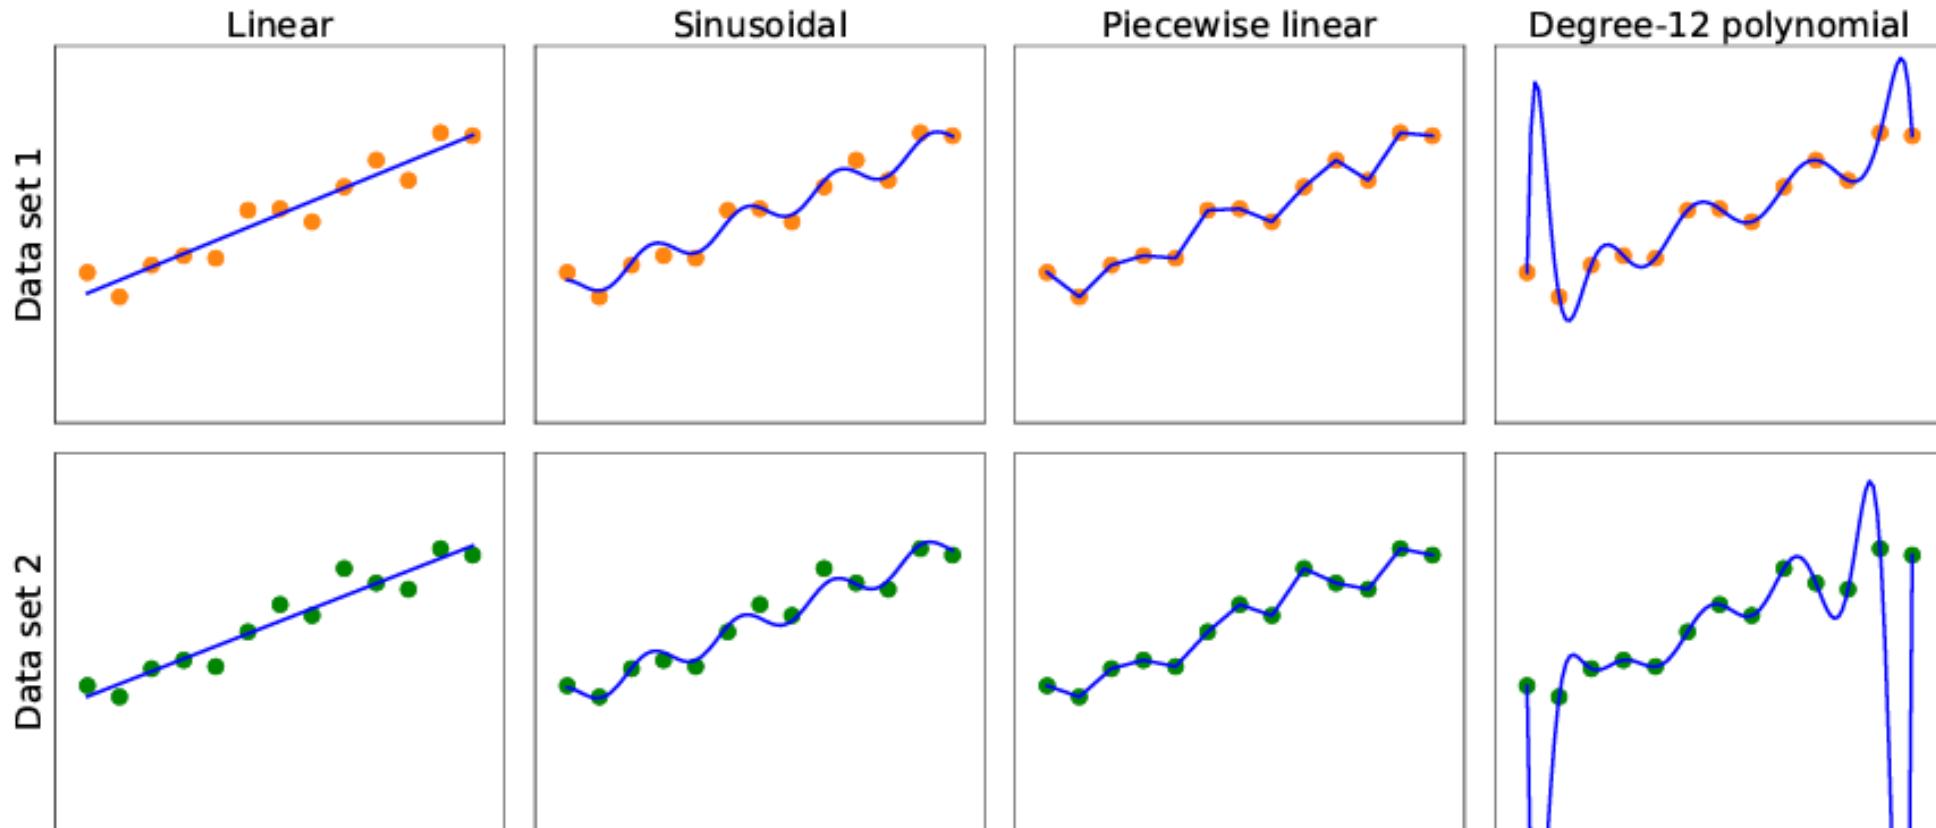
$$(x_1, y_1), (x_2, y_2), \dots (x_N, y_N),$$

where each pair was generated by an unknown function $y = f(x)$,
discover a function h that approximates the true function f .

- The function h is called a **hypothesis** or a **model** of the data

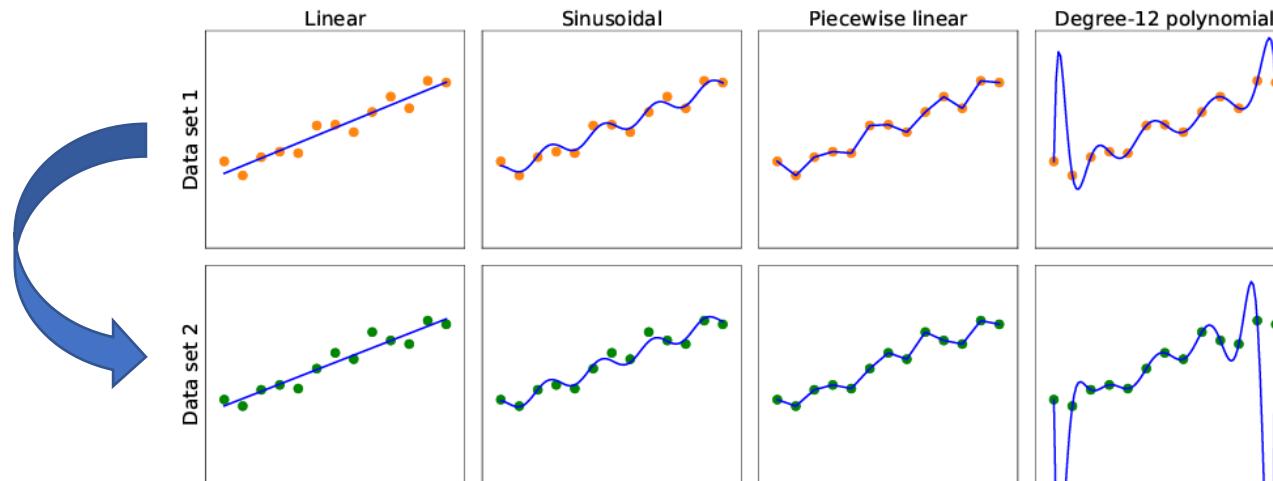


Supervised Learning



Supervised Learning - Generalization

- Performance measures
 - Performance on the training set
 - Performance on test set
- **Generalization**
 - h generalizes well if it accurately predicts the outputs of the **test set**



Supervised Learning - Bias & Variance

- One way to analyze **hypothesis spaces** is by the **bias** they impose (regardless of the training data set) and the **variance** they produce (from one training set to another).
- **Bias**
 - The tendency of a predictive hypothesis to deviate from the expected value when averaged over different training sets
- **Variance**
 - The amount of change in the hypothesis due to fluctuation in the training data
- **Bias-variance tradeoff**
 - more complex, low-bias hypotheses that fit the training data well
 - simpler, low-variance hypotheses that may generalize better



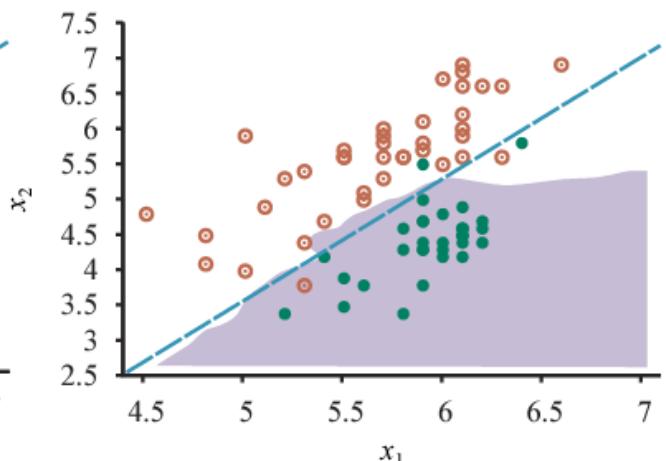
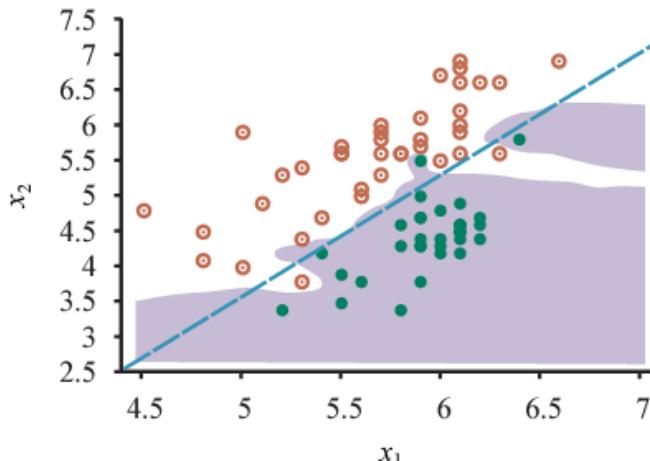
Supervised Learning - Underfitting & Overfitting

- **Underfitting**

- The algorithm fails to find a pattern in the data
- Model is too simple

- **Overfitting**

- Fits too much to the **particular data set** it is trained on
- Performs poorly on unseen data (train-test accuracy)
- Model is too complex
- Amount of data is not enough
- Data is not diverse enough



Decision Trees

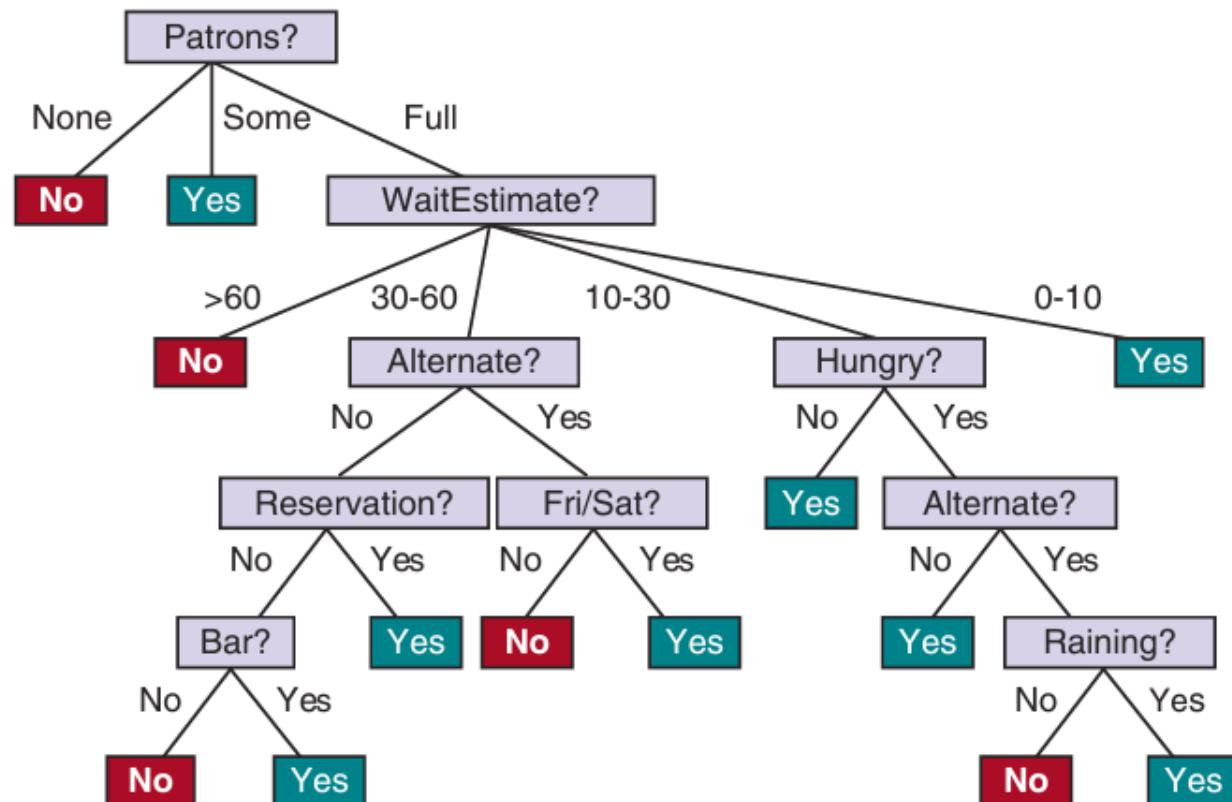
- Maps a vector of **attribute values** to a single output “**decision**”
- Sequence of tests
 - Start from root
 - Explore branches
 - Until a leaf is reached
- **Node = test of input attribute value**
- Input and output values can be **discrete** or **continuous**
- Simple example: **Boolean decision tree**
- A Boolean decision tree is equivalent to a logical statement of the form:
$$Output \Leftrightarrow (Path_1 \vee Path_2 \vee \dots),$$

$$(A_m = v_x \wedge A_n = v_y \wedge \dots)$$
- Where
the root to a true leaf
are attribute-value tests corresponding to a path from



Decision Trees

- Example: Waiting in restaurant



Decision Trees

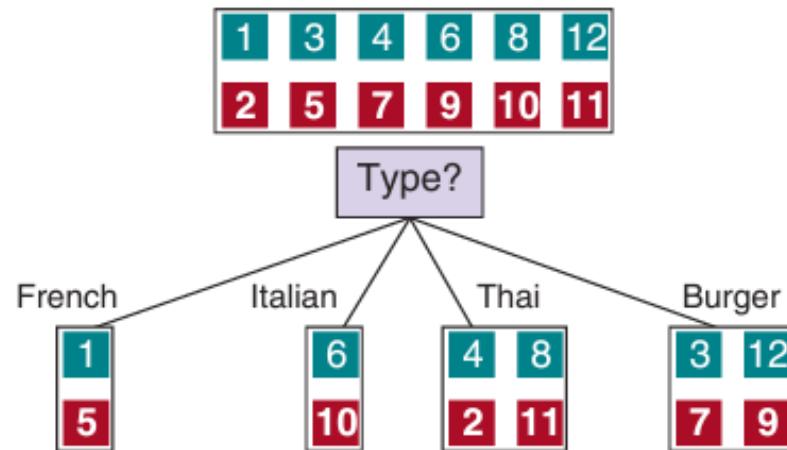
- Training dataset

Example	Input Attributes											Output WillWait
	Alt	Bar	Fri	Hun	Pat	Price	Rain	Res	Type	Est		
x ₁	Yes	No	No	Yes	Some	\$\$\$	No	Yes	French	0–10	y ₁ = Yes	
x ₂	Yes	No	No	Yes	Full	\$	No	No	Thai	30–60	y ₂ = No	
x ₃	No	Yes	No	No	Some	\$	No	No	Burger	0–10	y ₃ = Yes	
x ₄	Yes	No	Yes	Yes	Full	\$	Yes	No	Thai	10–30	y ₄ = Yes	
x ₅	Yes	No	Yes	No	Full	\$\$\$	No	Yes	French	>60	y ₅ = No	
x ₆	No	Yes	No	Yes	Some	\$\$	Yes	Yes	Italian	0–10	y ₆ = Yes	
x ₇	No	Yes	No	No	None	\$	Yes	No	Burger	0–10	y ₇ = No	
x ₈	No	No	No	Yes	Some	\$\$	Yes	Yes	Thai	0–10	y ₈ = Yes	
x ₉	No	Yes	Yes	No	Full	\$	Yes	No	Burger	>60	y ₉ = No	
x ₁₀	Yes	Yes	Yes	Yes	Full	\$\$\$	No	Yes	Italian	10–30	y ₁₀ = No	
x ₁₁	No	No	No	No	None	\$	No	No	Thai	0–10	y ₁₁ = No	
x ₁₂	Yes	Yes	Yes	Yes	Full	\$	No	No	Burger	30–60	y ₁₂ = Yes	

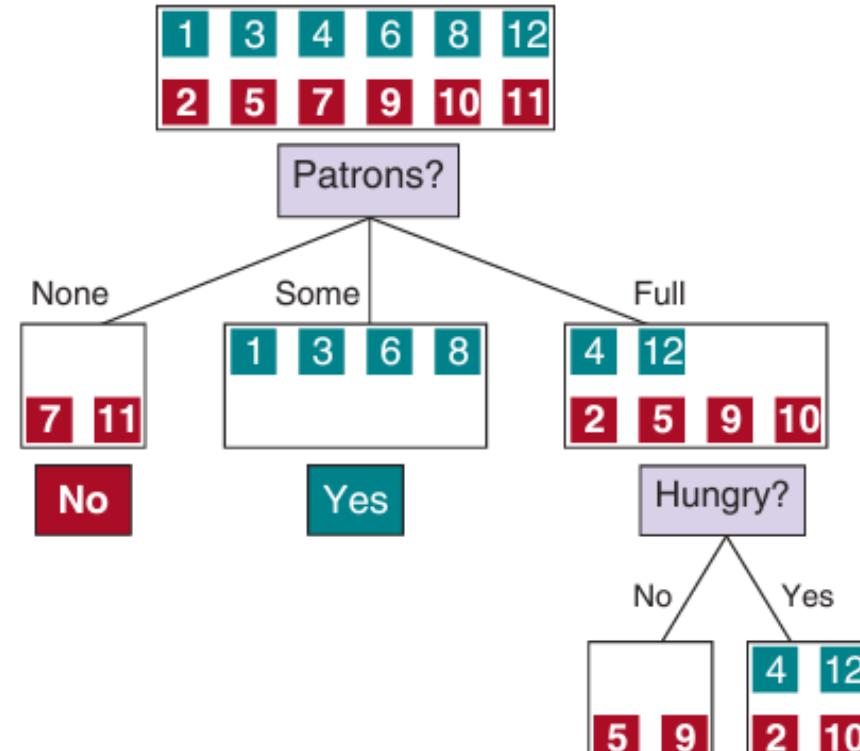


Decision Trees

- Find most important features and solve sub-problems sequentially



(a)



(b)



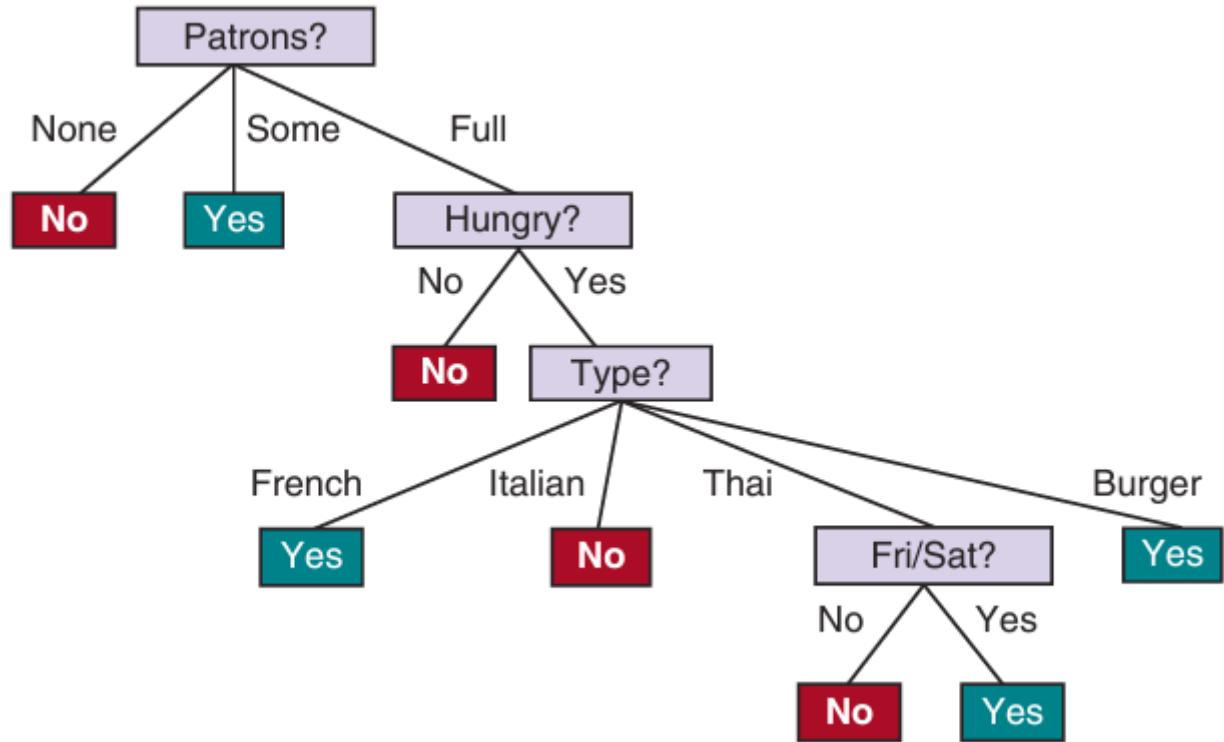
Decision Trees

```
function LEARN-DECISION-TREE(examples, attributes, parent-examples) returns a tree
  if examples is empty then return PLURALITY-VALUE(parent-examples)
  else if all examples have the same classification then return the classification
  else if attributes is empty then return PLURALITY-VALUE(examples)
  else
     $A \leftarrow \operatorname{argmax}_{a \in \text{attributes}} \text{IMPORTANCE}(a, \text{examples})$ 
    tree  $\leftarrow$  a new decision tree with root test A
    for each value v of A do
       $exs \leftarrow \{e : e \in \text{examples} \text{ and } e.A = v\}$ 
      subtree  $\leftarrow$  LEARN-DECISION-TREE(exs, attributes - A, examples)
      add a branch to tree with label (A = v) and subtree subtree
  return tree
```



Decision Trees

- Final decision tree has short paths
- Simpler than the original tree
- Fitted to the training set
- **BUT** The learning algorithm looks at the examples, not at the correct function
- Selecting most important attributes based on **information gain** defined by **expected reduction in entropy** of the output variable



Decision Trees

- **Generalization**
 - If we **increase** the **number of attributes**, overfitting is **more likely**
 - If we **increase** the **number of training samples**, overfitting is **less likely**
- **Pruning:** improving generalization for decision trees
 - Eliminate nodes that are not clearly relevant
 - Look at a test node that has only leaf nodes as descendants
 - If the test appears to be irrelevant - eliminate and replace it with a leaf node
 - Use significance test to measure low information gain threshold



Decision Trees

- **Pros**

- Ease of understanding
- Scalability to large data sets
- Versatility in handling discrete and continuous inputs
- Performing classification and regression

- **Cons**

- Suboptimal accuracy (largely due to the greedy search)
- If trees are very deep, then getting a prediction for a new example can be expensive
- Decision trees are unstable – adding just one new example can change the entire tree



Linear Regression and Classification

- **Univariate Linear Regression** = “fitting a straight line”

*fit model on (x, y) training examples so that $y = w_1x + w_0$
where w_0 and w_1 are learned weights*

- The learned linear function

$$h_{\mathbf{w}}(x) = w_1x + w_0$$

- The task of finding this linear function based on training data is called **linear regression**

- Find values of w_0 , w_1 that minimizes the empirical loss (e.g. L2 loss)

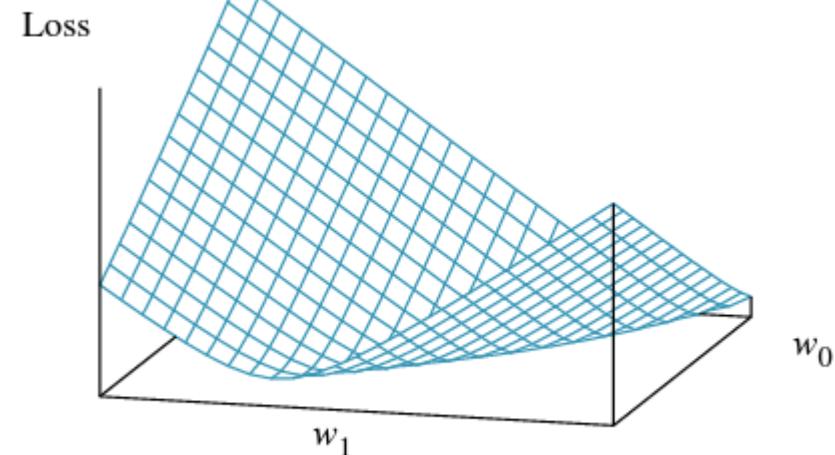
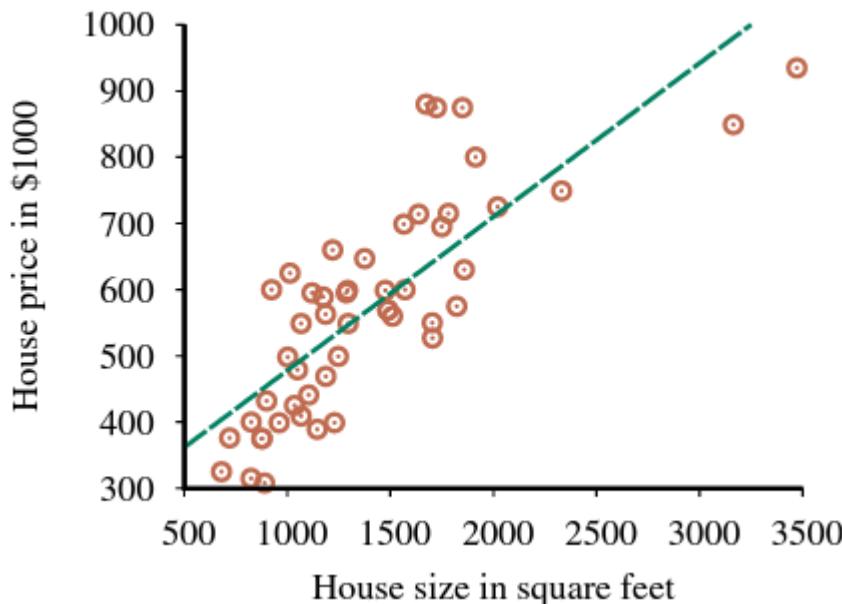
$$\text{Loss}(h_{\mathbf{w}}) = \sum_{j=1}^N L_2(y_j, h_{\mathbf{w}}(x_j)) = \sum_{j=1}^N (y_j - h_{\mathbf{w}}(x_j))^2 = \sum_{j=1}^N (y_j - (w_1x_j + w_0))^2$$



Linear Regression and Classification

- We can find the function that minimizes the loss using partial derivatives

$$\frac{\partial}{\partial w_0} \sum_{j=1}^N (y_j - (w_1 x_j + w_0))^2 = 0 \text{ and } \frac{\partial}{\partial w_1} \sum_{j=1}^N (y_j - (w_1 x_j + w_0))^2 = 0$$



Linear Regression and Classification

- For more complex cases we need to introduce **gradient descent**
 - Algorithm for finding a point in the **weight space** that **minimizes the loss**
 - Similar to the hill climbing but here we are minimizing the loss, not maximizing the gain
- **Gradient Descent**
 - Select random starting point in weight space
 - Compute an estimate of the gradient
 - Move a small amount in the steepest downhill direction
 - Repeat until we converge on a point with minimum loss
- Parameters
 - Learning rate (α) - determines the step size for the descent

```
w ← any point in the parameter space  
while not converged do  
    for each  $w_i$  in w do  
         $w_i \leftarrow w_i - \alpha \frac{\partial}{\partial w_i} Loss(\mathbf{w})$ 
```



Linear Regression and Classification

- For the univariate case the loss was quadratic - the partial derivative is linear

- We apply the chain rule

$$\frac{\partial}{\partial w_i} \text{Loss}(\mathbf{w}) = \frac{\partial}{\partial w_i} (y - h_{\mathbf{w}}(x))^2 = 2(y - h_{\mathbf{w}}(x)) \times \frac{\partial}{\partial w_i} (y - h_{\mathbf{w}}(x))$$

$$= 2(y - h_{\mathbf{w}}(x)) \times \frac{\partial}{\partial w_i} (y - (w_1 x + w_0)).$$

$$\frac{\partial}{\partial w_0} \text{Loss}(\mathbf{w}) = -2(y - h_{\mathbf{w}}(x)); \quad \frac{\partial}{\partial w_1} \text{Loss}(\mathbf{w}) = -2(y - h_{\mathbf{w}}(x)) \times x$$

- Create the update function with the selected learning rate

$$w_0 \leftarrow w_0 + \alpha (y - h_{\mathbf{w}}(x)); \quad w_1 \leftarrow w_1 + \alpha (y - h_{\mathbf{w}}(x)) \times x$$

Single example

$$w_0 \leftarrow w_0 + \alpha \sum_j (y_j - h_{\mathbf{w}}(x_j)); \quad w_1 \leftarrow w_1 + \alpha \sum_j (y_j - h_{\mathbf{w}}(x_j)) \times x_j$$

Batch

